


```
In [ ]: x_train3=(x_train2>0)*1
x_test3 =(x_test2>0)*1
```


a) visualize data

```
In [ ]: import warnings
warnings.filterwarnings("ignore")
```

Standardized data diagram

```
In [ ]: data1=pd.DataFrame(x_train1)
data1["label"]=y_train
radviz(data1, 'label')
```


Out[]: <AxesSubplot:>



Log-transformed data diagram

```
In [ ]: data2=pd.DataFrame(x_train2)
data2["label"]=y_train
radviz(data2, 'label')
```

Out[]: <AxesSubplot:>




Discretized data diagram

```
In [ ]: data3=pd.DataFrame(x_train3)
data3["label"]=y_train
```

```
radviz(data3, 'label')
```

Out[]: <AxesSubplot:>



b) fit logistic model

i) fit on standardized data

```
In [ ]: # fit on standardized data
X=sm.add_constant(x_train1)
y=y_train
logit_data1=sm.Logit(y,X).fit()

# error rate on train set
logit_train_pred1=(logit_data1.predict(train)>0.5).astype(int)
logit_train_error1=sum(logit_train_pred1!=y_train)/len(y_train)
print(logit_train_error1)

# error rate on test set
logit_test_pred1=(logit_data1.predict(test)>0.5).astype(int)
logit_test_error1=sum(logit_test_pred1!=y_test)/len(y_test)
print(logit_test_error1)

# Significant features: p<0.05 is defined as significant
logit_data1.pvalues[logit_data1.pvalues<0.05].index[1:]

Optimization terminated successfully.
    Current function value: 0.188692
    Iterations 15
0.5500489077274209
0.5345501955671447
Out[ ]: Index(['x4', 'x5', 'x7', 'x8', 'x9', 'x11', 'x16', 'x17', 'x19', 'x20', 'x21',
              'x23', 'x25', 'x27', 'x42', 'x44', 'x45', 'x46', 'x47', 'x49', 'x52',
              'x53', 'x55', 'x56', 'x57'],
              dtype='object')
```

ii) fit on log data

```
In [ ]: # fit on log data
X=sm.add_constant(x_train2)
y=y_train
logit_data2=sm.Logit(y,X).fit()

# error rate on train set
```

```

logit_train_pred2=(logit_data2.predict(train)>0.5).astype(int)
logit_train_error2=sum(logit_train_pred2!=y_train)/len(y_train)
print(logit_train_error2)

# error rate on test set
logit_test_pred2=(logit_data2.predict(test)>0.5).astype(int)
logit_test_error2=sum(logit_test_pred2!=y_test)/len(y_test)
print(logit_test_error2)

# Significant features: p<0.05 is defined as significant
logit_data1.pvalues[logit_data2.pvalues<0.05].index[1:]

```

Optimization terminated successfully.

Current function value: 0.154490

Iterations 14

0.527551353113792

0.5176010430247718

Out[]: Index(['x5', 'x7', 'x8', 'x11', 'x13', 'x15', 'x16', 'x17', 'x18', 'x20',
'x21', 'x23', 'x24', 'x25', 'x27', 'x28', 'x35', 'x37', 'x42', 'x43',
'x44', 'x45', 'x46', 'x52', 'x53', 'x57'],
dtype='object')

iii) fit on discretize data

```

In [ ]: # fit on discretize data
X=sm.add_constant(x_train3)
y=y_train
logit_data3=sm.Logit(y,X).fit()

# error rate on train set
logit_train_pred3=(logit_data3.predict(train)>0.5).astype(int)
logit_train_error3=sum(logit_train_pred3!=y_train)/len(y_train)
print(logit_train_error3)

# error rate on test set
logit_test_pred3=(logit_data3.predict(test)>0.5).astype(int)
logit_test_error3=sum(logit_test_pred3!=y_test)/len(y_test)
print(logit_test_error3)

# Significant features: p<0.05 is defined as significant
logit_data1.pvalues[logit_data3.pvalues<0.05].index[1:]

```

Optimization terminated successfully.

Current function value: 0.173849

Iterations 12

0.46723182262797525

0.4485006518904824

Out[]: Index(['x5', 'x7', 'x8', 'x10', 'x13', 'x15', 'x16', 'x17', 'x18', 'x20',
'x21', 'x22', 'x23', 'x24', 'x25', 'x26', 'x27', 'x28', 'x31', 'x33',
'x35', 'x37', 'x42', 'x44', 'x45', 'x46', 'x52', 'x53', 'x54', 'x55',
'x57'],
dtype='object')

c) fit LDA/QDA

LDA method

```

In [ ]: # fit on standardized data
lda = LinearDiscriminantAnalysis()
lda.fit(x_train1, y)

```

```

# error rate on train set
lda_train_pred1=lda.predict(x_train1)
lda_train_error1=sum(lda_train_pred1!=y_train)/len(y_train)
print(lda_train_error1)

# error rate on test set
lda_test_pred1=lda.predict(x_test1)
lda_test_error1=sum(lda_test_pred1!=y_test)/len(y_test)
print(lda_test_error1)

```

0.10172807303553962

0.10299869621903521

```

In [ ]: # fit on log data
lda = LinearDiscriminantAnalysis()
lda.fit(x_train2, y)

# error rate on train set
lda_train_pred2=lda.predict(x_train2)
lda_train_error2=sum(lda_train_pred2!=y_train)/len(y_train)
print(lda_train_error2)

# error rate on test set
lda_test_pred2=lda.predict(x_test2)
lda_test_error2=sum(lda_test_pred2!=y_test)/len(y_test)
print(lda_test_error2)

```

0.07042712748614281

0.07627118644067797

QDA method

```

In [ ]: # fit on standardized data
qda = QuadraticDiscriminantAnalysis()
qda.fit(x_train1, y)

# error rate on train set
qda_train_pred1=qda.predict(x_train1)
qda_train_error1=sum(qda_train_pred1!=y_train)/len(y_train)
print(qda_train_error1)

# error rate on test set
qda_test_pred1=qda.predict(x_test1)
qda_test_error1=sum(qda_test_pred1!=y_test)/len(y_test)
print(qda_test_error1)

```

0.17867623084447343

0.17470664928292046

```

In [ ]: # fit on log data
qda = QuadraticDiscriminantAnalysis()
qda.fit(x_train2, y)

# error rate on train set
qda_train_pred2=qda.predict(x_train2)
qda_train_error2=sum(qda_train_pred2!=y_train)/len(y_train)
print(qda_train_error2)

# error rate on test set
qda_test_pred2=qda.predict(x_test2)
qda_test_error2=sum(qda_test_pred2!=y_test)/len(y_test)
print(qda_test_error2)

```

0.15194000652103032
0.14993481095176012

d) fit svm

linear support vector machine

```
In [ ]: # fit on standardized data
Svm =svm.SVC(kernel='linear')
Svm.fit(x_train1, y)

# error rate on train set
svm_train_pred1=Svm.predict(x_train1)
svm_train_error1=sum(svm_train_pred1!=y_train)/len(y_train)
print(svm_train_error1)

# error rate on test set
svm_test_pred1=Svm.predict(x_test1)
svm_test_error1=sum(svm_test_pred1!=y_test)/len(y_test)
print(svm_test_error1)

0.06488425171177047
0.07170795306388526
```

```
In [ ]: # fit on standardized data
Svm =svm.SVC(kernel='linear')
Svm.fit(x_train2, y)

# error rate on train set
svm_train_pred2=Svm.predict(x_train2)
svm_train_error2=sum(svm_train_pred2!=y_train)/len(y_train)
print(svm_train_error2)

# error rate on test set
svm_test_pred2=Svm.predict(x_test2)
svm_test_error2=sum(svm_test_pred2!=y_test)/len(y_test)
print(svm_test_error2)

0.056080860776002606
0.06258148631029987
```

```
In [ ]: # fit on discretize data
Svm =svm.SVC(kernel='linear')
Svm.fit(x_train3, y)

# error rate on train set
svm_train_pred3=Svm.predict(x_train3)
svm_train_error3=sum(svm_train_pred3!=y_train)/len(y_train)
print(svm_train_error3)

# error rate on test set
svm_test_pred3=Svm.predict(x_test3)
svm_test_error3=sum(svm_test_pred3!=y_test)/len(y_test)
print(svm_test_error3)

0.06423214867949135
0.07301173402868318
```

non-linear support vector machine

```
In [ ]: # fit on standardized data
```

```
Svm =svm.SVC(kernel='rbf')
Svm.fit(x_train1, y)

# error rate on train set
n_svm_train_pred1=Svm.predict(x_train1)
n_svm_train_error1=sum(svm_train_pred1!=y_train)/len(y_train)
print(n_svm_train_error1)

# error rate on test set
n_svm_test_pred1=Svm.predict(x_test1)
n_svm_test_error1=sum(n_svm_test_pred1!=y_test)/len(y_test)
print(n_svm_test_error1)

0.06488425171177047
0.06453715775749674
```

```
In [ ]: # fit on log data
Svm =svm.SVC(kernel='rbf')
Svm.fit(x_train2, y)

# error rate on train set
n_svm_train_pred2=Svm.predict(x_train2)
n_svm_train_error2=sum(svm_train_pred2!=y_train)/len(y_train)
print(n_svm_train_error2)

# error rate on test set
n_svm_test_pred2=Svm.predict(x_test2)
n_svm_test_error2=sum(n_svm_test_pred2!=y_test)/len(y_test)
print(n_svm_test_error2)

0.056080860776002606
0.04954367666232073
```

```
In [ ]: # fit on discretize data
Svm =svm.SVC(kernel='rbf')
Svm.fit(x_train3, y)

# error rate on train set
n_svm_train_pred3=Svm.predict(x_train3)
n_svm_train_error3=sum(svm_train_pred3!=y_train)/len(y_train)
print(n_svm_train_error3)

# error rate on test set
n_svm_test_pred3=Svm.predict(x_test3)
n_svm_test_error3=sum(n_svm_test_pred3!=y_test)/len(y_test)
print(n_svm_test_error3)

0.06423214867949135
0.05345501955671447
```

e) decision tree classifier

```
In [ ]: # fit on standardized data
dc = RandomForestClassifier(max_depth=6)
dc.fit(x_train1, y_train)

# error rate on train set
dc_train_pred1=dc.predict(x_train1)
dc_train_error1=sum(dc_train_pred1!=y_train)/len(y_train)
print(dc_train_error1)

# error rate on test set
dc_test_pred1=dc.predict(x_test1)
```

```
dc_test_error1=sum(dc_test_pred1!=y_test)/len(y_test)
print(dc_test_error1)
```

```
0.051516139550048905
0.05410691003911343
```

```
In [ ]: # fit on log data
dc = RandomForestClassifier(max_depth=6)
dc.fit(x_train2, y_train)

# error rate on train set
dc_train_pred2=dc.predict(x_train2)
dc_train_error2=sum(dc_train_pred2!=y_train)/len(y_train)
print(dc_train_error2)

# error rate on test set
dc_test_pred2=dc.predict(x_test2)
dc_test_error2=sum(dc_test_pred2!=y_test)/len(y_test)
print(dc_test_error2)
```

```
0.05086403651776981
0.052803129074315516
```

```
In [ ]: # fit on log data
dc = RandomForestClassifier(max_depth=6)
dc.fit(x_train3, y_train)

# error rate on train set
dc_train_pred3=dc.predict(x_train3)
dc_train_error3=sum(dc_train_pred3!=y_train)/len(y_train)
print(dc_train_error3)

# error rate on test set
dc_test_pred3=dc.predict(x_test3)
dc_test_error3=sum(dc_test_pred3!=y_test)/len(y_test)
print(dc_test_error3)
```

```
0.07042712748614281
0.07496740547588006
```

Report classification errors

```
In [ ]: Methods=["Logit", "LDA", "QDA", "SVM-linear", "SVM-nonlinear", "Decision Tree"]
Methods = [val for val in Methods for _ in range(3)]
Data=["standardized", "log-transformed", "Discretized"]*6
Train_error_rate=[logit_train_error1, logit_train_error2, logit_train_error3,
                  lda_train_error1,   lda_train_error2,   "----",
                  qda_train_error1,   qda_train_error2,   "----",
                  svm_train_error1,   svm_train_error2,   svm_train_error3,
                  n_svm_train_error1, n_svm_train_error2, n_svm_train_error3,
                  dc_train_error1,   dc_train_error2,   dc_train_error3]
Test_error_rate=[logit_test_error1, logit_test_error2, logit_test_error3,
                 lda_test_error1,   lda_test_error2,   "----",
                 qda_test_error1,   qda_test_error2,   "----",
                 svm_test_error1,   svm_test_error2,   svm_test_error3,
                 n_svm_test_error1, n_svm_test_error2, n_svm_test_error3,
                 dc_test_error1,   dc_test_error2,   dc_test_error3]
result=pd.DataFrame({"Methods":Methods,
                    "Data":Data,
                    "Train_error_rate":Train_error_rate,
                    "Test_error_rate":Test_error_rate})
result
```


Out[]:

	Methods	Data	Train_error_rate	Test_error_rate
0	Logit	standardized	0.550049	0.53455
1	Logit	log-transformed	0.527551	0.517601
2	Logit	Discretized	0.467232	0.448501
3	LDA	standardized	0.101728	0.102999
4	LDA	log-transformed	0.070427	0.076271
5	LDA	Discretized	---	---
6	QDA	standardized	0.178676	0.174707
7	QDA	log-transformed	0.15194	0.149935
8	QDA	Discretized	---	---
9	SVM-linear	standardized	0.064884	0.071708
10	SVM-linear	log-transformed	0.056081	0.062581
11	SVM-linear	Discretized	0.064232	0.073012
12	SVM-nonlinear	standardized	0.064884	0.064537
13	SVM-nonlinear	log-transformed	0.056081	0.049544
14	SVM-nonlinear	Discretized	0.064232	0.053455
15	Decision Tree	standardized	0.051516	0.054107
16	Decision Tree	log-transformed	0.050864	0.052803
17	Decision Tree	Discretized	0.070427	0.074967

In the table above, we see the nonlinear support vector machine on log-transformed data performs best, which has the lowest test error rate and also fits well on training data. The worst model is logit model, the test error rate is approximately close to 0.5.

Finally, we choose nonlinear support vector machine on log-transformed data as our best model. Then, we try different soft margin param C, and find the optimal option to make the test error rate as small as possible.

```
In [ ]: #Find best C
scores_list = []
C_list = 10*np.linspace(-2,5,100)
for C_val in C_list:
    model = svm.SVC(kernel='rbf', C=C_val)
    model.fit(x_train2,y_train)
    pred=model.predict(x_test2)
    test_error=sum(pred!=y_test)/len(y_test)
    scores_list.append(test_error)

plt.plot(C_list, scores_list, color = 'blue', marker = '.', markersize = 8,
         markeredgecolor = 'black', markerfacecolor = 'black',label = 'Score')
plt.title('Accuracy Score vs C')
plt.show()
```

```

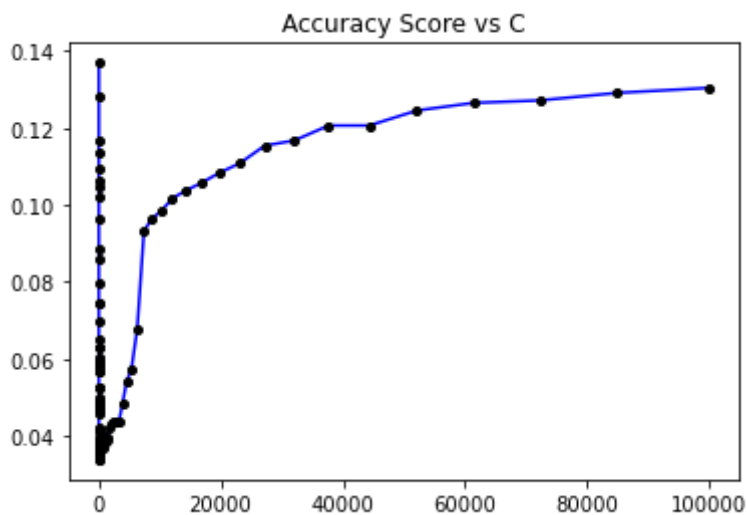
index = np.argmin(np.array(scores_list))
C_best = C_list[index]
print('The best C is ', C_best)

#Find best gamma
scores_list = []
gamma_list = 10*np.linspace(-2, 5, 100)
for g_val in gamma_list:
    model = svm.SVC(kernel='rbf', C=C_best, gamma=g_val)
    model.fit(x_train2, y_train)
    pred = model.predict(x_test2)
    test_error = sum(pred != y_test)/len(y_test)
    scores_list.append(test_error)

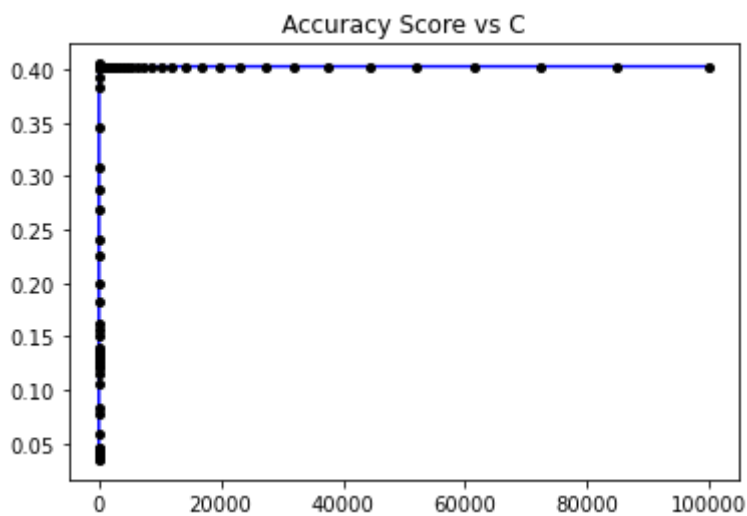
plt.plot(gamma_list, scores_list, color='blue', marker='.', markersize=8,
         markeredgecolor='black', markerfacecolor='black', label='Score')
plt.title('Accuracy Score vs C')
plt.show()

index = np.argmin(np.array(scores_list))
g_best = gamma_list[index]
print('The best gamma is ', g_best)

```



The best C is 77.4263682681127



The best gamma is 0.05994842503189409

The best hyperparameter C for SVM model is 77.42. And best gamma is 0.05995.

```

In [ ]: model = svm.SVC(kernel='rbf', C=C_best, gamma=g_best)
        model.fit(x_train2, y_train)
        pred=model.predict(x_test2)

```

```
test_error=sum(pred!=y_test)/len(y_test)
test_error
```

Out[]: 0.03455019556714472

Now the test error rate has been decreased to 3.46%.It does a good job!

Everyone contributed equally.